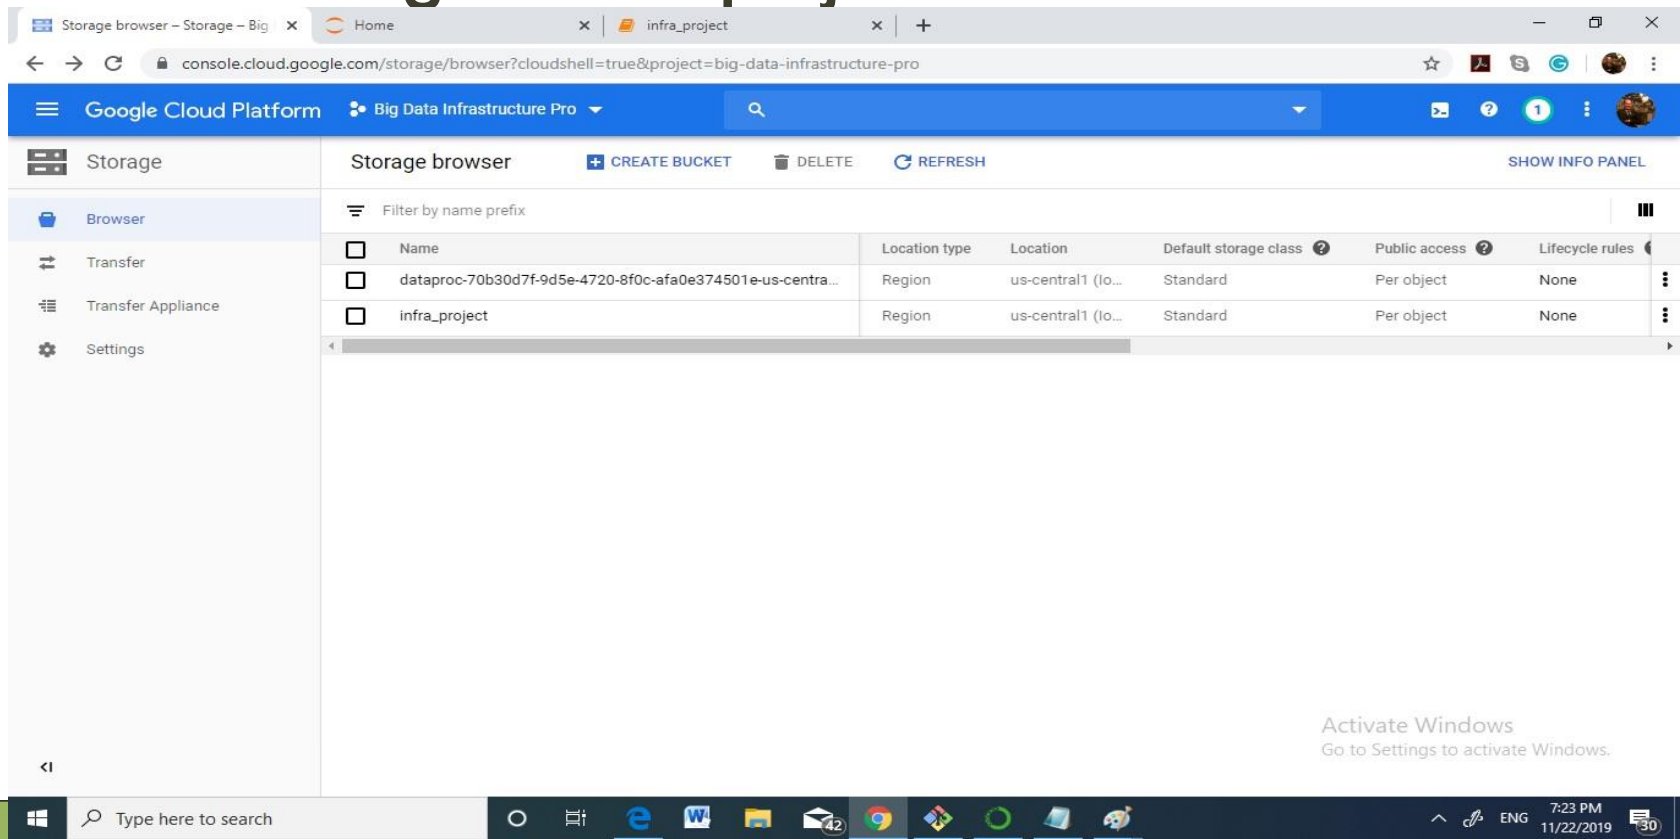


Monthly prediction for Bixi rides in Montreal

The final project for big data
infrastructure course by
Hamid Reza Taremian

Setting up Google cloud platform

- 1. First we need to set up a bucket to have a storage for our project



The screenshot displays the Google Cloud Platform Storage browser interface. The left sidebar shows the 'Storage' section with options for 'Browser', 'Transfer', 'Transfer Appliance', and 'Settings'. The main area, titled 'Storage browser', includes buttons for 'CREATE BUCKET', 'DELETE', and 'REFRESH', along with a 'SHOW INFO PANEL' link. A search bar labeled 'Filter by name prefix' is present. Below it, a table lists storage buckets:

<input type="checkbox"/>	Name	Location type	Location	Default storage class	Public access	Lifecycle rules
<input type="checkbox"/>	dataproc-70b30d7f-9d5e-4720-8f0c-afa0e374501e-us-centra...	Region	us-central1 (lo...	Standard	Per object	None
<input type="checkbox"/>	infra_project	Region	us-central1 (lo...	Standard	Per object	None

An 'Activate Windows' watermark is visible in the bottom right corner of the interface.

- 2. Then we need to make a cluster for our project using command shell which will allow jupyter notebook using the following command:

- `gcloud beta dataproc clusters create infra_project`
`\ --optional-components=ANACONDA,JUPYTER`
`\ --image-version=1.3`
`\ --enable-component-gateway`
`\ --bucket infra_project`
`\ --project big-data-infrastructure-pro`

Google Cloud Platform Big Data Infrastructure Pro

Dataprocc

Clusters

CREATE CLUSTER REFRESH DELETE REGIONS SHOW INFO PANEL

Search clusters, press Enter

Name	Region	Zone	Total worker nodes	Scheduled deletion	Cloud Storage staging bucket	Created	Status
projcet	us-central1	us-central1-f	2	Off	infra_project	Nov 22, 2019, 6:15:14 PM	Running

(big-data-infrastructure-pro) x +

```
Welcome to Cloud Shell! Type "help" to get started.
Your Cloud Platform project in this session is set to big-data-infrastructure-pro.
Use "gcloud config set project [PROJECT ID]" to change to a different project.
hamidreza_tareimian@cloudshell:~ (big-data-infrastructure-pro)$ gcloud beta dataprocc clusters create projcet --optional-components=ANACONDA,JUPYTER --image-version=1.3 --enable-component-gateway --bucket infra_project --project big-data-infrastructure-pro --region us-central1
```

Activate Windows
Go to Settings to activate Windows.

Some Analysis

- we can see how many rides were done from each station and link the station geographical information for further usage.

3. Now we can have access to jupyter notebook to write our code

Bucket det... x | G Namenode x | G Namenode x | Home x | infra_project x | Untitled x | Markdown x | Infra_project x | git - How d x | + - □ x

← → ↻ e4cyqydsnbtzobm3iyuiaxjr4-dot-us-central1.dataproc.googleusercontent.com/gateway/default/jupyter/notebooks/infra_project.ipynb# ☆ 📄 S G 🌐

jupyter infra_project Last Checkpoint: 43 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted | PySpark ○

📄 + 🔍 📄 ⬆ ⬇ ▶ Run ■ ↺ ▶ Markdown ▾ 🗨

```
## Starting and configuring our spark master and session
```

```
In [1]: from pyspark.sql.types import StructType
format = StructType().add("Code", "string").add("Name", "string").add("latitude", "float").add("longitude", "float")
```

```
In [2]: from pyspark.context import SparkContext
from pyspark.sql.session import SparkSession
#import geopandas
import datetime
```

```
In [3]: #sc = SparkContext()

spark = SparkSession.builder \
    .master('spark://master:7077') \
    .appName("cebd1261") \
    .getOrCreate()
```

```
### *reading the data related to the stations like:
#### *station code
#### * latitude and longitude of the station
#### * the name of the station
```

infra_project (1).ipynb

Activate Windows
Go to Settings to activate Windows.
Show all x

Type here to search

9:21 PM 11/22/2019

The python+spark code

- our data consists of one CSV file for the information related to the each station and CSV files for monthly rides. so first we upload our files into the bucket that we made before.

Bucket details - Big Data Infrastru x Home x infra_project x +

console.cloud.google.com/storage/browser/infra_project/?cloudshell=true&project=big-data-infrastructure-pro

Google Cloud Platform Big Data Infrastructure Pro

Storage

Browser

Transfer

Transfer Appliance

Settings

Bucket details

EDIT BUCKET

REFRESH BUCKET

infra_project

Objects

Overview

Permissions

Bucket Lock

Upload files

Upload folder

Create folder

Manage holds

Delete

Filter by prefix...

Buckets / infra_project

					UTC-5				
<input type="checkbox"/>	OD_2019-08.csv	57.14 MB	application/vnd.ms-excel	Standard	11/22/19, 6:56:38 PM UTC-5	Not public	Google-managed key	-	None
<input type="checkbox"/>	OD_2019-09.csv	47.56 MB	application/vnd.ms-excel	Standard	11/22/19, 6:56:47 PM UTC-5	Not public	Google-managed key	-	None
<input type="checkbox"/>	OD_2019-10.csv	33.68 MB	application/vnd.ms-excel	Standard	11/22/19, 6:50:02 PM UTC-5	Not public	Google-managed key	-	None
<input type="checkbox"/>	OD_2019.csv/	-	Folder	-	-	Per object	-	-	-
<input type="checkbox"/>	google-cloud-dataproc-metainfo/	-	Folder	-	-	Per object	-	-	-
<input type="checkbox"/>	notebooks/	-	Folder	-	-	Per object	-	-	-

Activate Windows

Go to Settings to activate Windows.

Type here to search

ENG 7:22 PM 11/22/2019


```
In [7]: station_cnt = df_2019_10.groupby('start_station_code').agg({'start_station_code': 'count'})
station_cnt.show()
```

```
+-----+-----+
|start_station_code|count(start_station_code)|
+-----+-----+
|          6194|          2301|
|          6731|           271|
|          6240|           383|
|          6248|          3293|
|          6366|           537|
|          6903|          2108|
|          7056|           301|
|          6081|           358|
|          6227|          2535|
|          7054|           696|
|          6380|           547|
|          6106|           520|
|          6732|           608|
|          6143|          1911|
|          6252|          1152|
|          6402|           722|
|          7014|          1635|
|          7013|           375|
|          6033|           424|
|          7093|           42|
+-----+-----+
only showing top 20 rows
```

Activate Windows

- we can read the data for stations and each month separately.

```
In [4]: df_station = spark.read.csv('gs://infra_project/Stations_2019.csv', schema=format, header="true")
df_station.show(5)
```

```
+-----+-----+-----+-----+
| Code|          Name| latitude|longitude|
+-----+-----+-----+-----+
|10002|Métro Charlevoix ...| 45.47823| -73.56965|
| 4000|Jeanne-d'Arc / On...| 45.5496| -73.54188|
| 4001|Graham / Brookfield|45.520073|-73.629776|
| 4002|Graham / Wicksteed|45.516937|-73.64048|
| 5002|St-Charles / Mont...|45.533684|-73.51526|
+-----+-----+-----+-----+
```

only showing top 5 rows

Code	Name	latitude	longitude	October count
6194	Métro Atwater (At...	45.489475	-73.584564	2301
6731	28e avenue / Rose...	45.564354	-73.57124	271
6240	Parc Kent (de Ken...	45.505722	-73.629456	383
6248	St-Dominique / Ra...	45.518593	-73.581566	3293
6366	Wilderton / Van ...	45.510143	-73.62475	537
6903	St-Dominique / Na...	45.516663	-73.57722	2108
7056	Bibliothèque de V...	45.44826	-73.57786	301
6081	Mackay / Ste-Cath...	45.49571	-73.57695	358
6227	de l'Esplanade / ...	45.521038	-73.59491	2535
7054	de la Côte St-Pau...	45.467667	-73.59392	696
6380	Parc J.-Arthur-Ch...	45.551582	-73.56191	547
6106	Papineau / René-L...	45.52114	-73.54926	520
6732	Fortune / Wellington	45.477924	-73.55904	608
6143	Rachel / de Brébeuf	45.52689	-73.57264	1911
6252	Mozart / St-Laurent	45.53318	-73.61544	1152
6402	Ste-Émilie / Sir-...	45.472668	-73.58539	722
7014	Métro Université ...	45.504276	-73.61797	1635
7013	Benny / Sherbrooke	45.464878	-73.626595	375
6033	16e avenue / Beau...	45.55828	-73.58316	424
7093	Laforest / Dudemaine	45.539806	-73.68726	42

only showing top 20 rows

- we can decide which rides happened during the same day and how many were more than one day

```
In [9]: def sameday(day1,day2):  
        if day1==day2:  
            same_day=1  
            print('yes')  
        else:  
            same_day=0  
        return(same_day)  
  
spark.udf.register("Check_same_day", sameday)
```

```
Out[9]: <function __main__.Check_same_day>
```

```
In [10]: from pyspark.sql.functions import month,year,dayofmonth,monotonically_increasing_id,udf,struct  
        from pyspark.sql.types import IntegerType  
  
q=df_2019_10.select(year(df_2019_10.start_date),month(df_2019_10.start_date),dayofmonth(df_2019_10.start_date),dayofmonth(df_2019_10.start_date))  
  
q=q.withColumnRenamed('month(start_date)','month').withColumnRenamed('year(start_date)','year').withColumnRenamed('dayofmonth(start_date)','dayofmonth')  
  
q = q.withColumn('index',monotonically_increasing_id())  
df_2019_10 = df_2019_10.withColumn('index',monotonically_increasing_id())  
  
df_2019_10=df_2019_10.join(q,on='index')
```

index	start_date	start_station_code	duration_sec	is_member	year	month	sday	end_day	same
0	2019-10-01 00:00:08	6174	199	1	2019	10	1	1	1
1	2019-10-01 00:01:13	6196	205	1	2019	10	1	1	1
2	2019-10-01 00:01:34	6033	212	1	2019	10	1	1	1
3	2019-10-01 00:02:33	6136	256	1	2019	10	1	1	1
4	2019-10-01 00:02:34	6204	104	1	2019	10	1	1	1
5	2019-10-01 00:02:49	6052	596	1	2019	10	1	1	1
6	2019-10-01 00:04:07	6149	503	1	2019	10	1	1	1
7	2019-10-01 00:04:25	7032	438	1	2019	10	1	1	1
8	2019-10-01 00:04:51	6118	262	1	2019	10	1	1	1
9	2019-10-01 00:05:01	6204	519	0	2019	10	1	1	1
10	2019-10-01 00:05:01	6095	540	1	2019	10	1	1	1
11	2019-10-01 00:05:52	6753	1111	1	2019	10	1	1	1
12	2019-10-01 00:06:23	6387	359	1	2019	10	1	1	1
13	2019-10-01 00:06:26	6181	168	1	2019	10	1	1	1
14	2019-10-01 00:06:41	6254	475	1	2019	10	1	1	1
15	2019-10-01 00:07:11	6432	1517	0	2019	10	1	1	1
16	2019-10-01 00:07:46	6021	375	1	2019	10	1	1	1
17	2019-10-01 00:09:19	6404	1176	1	2019	10	1	1	1
18	2019-10-01 00:09:51	6100	466	1	2019	10	1	1	1
19	2019-10-01 00:11:00	6184	654	1	2019	10	1	1	1

only showing top 20 rows


predictions

- The data at hand is time related so we have to run time series predictions on it. The step will be as follows:
 - extract the average time per ride and total number of rides per month for each station
 - put all these together to have the information for year 2019
 - run time series prediction

- we can extract the average time and number of rides per station using the code below:

```
In [13]: df_2019_10_agg=df_2019_10.groupby('year','month','sday').agg({'sday':'count','duration_sec':'mean'})
df_2019_10_agg=df_2019_10_agg.orderBy('sday')
df_2019_10_agg.show(31)
```

year	month	sday	count(sday)	avg(duration_sec)
2019	10	1	14017	636.348148676607
2019	10	2	26547	679.7920292311749
2019	10	3	24542	670.0463694890392
2019	10	4	23822	685.766392410377
2019	10	5	19613	773.0914189568143
2019	10	6	15544	688.2258749356665
2019	10	7	18289	664.0854612061895
2019	10	8	28226	722.1517395309289
2019	10	9	26730	702.8726898615788
2019	10	10	27761	713.6801988400994
2019	10	11	27242	733.204573819837
2019	10	12	18073	797.5678636640292
2019	10	13	17879	840.4145086414229
2019	10	14	13734	721.5892675112858
2019	10	15	24828	693.0045513130336
2019	10	16	22370	682.5941439427805
2019	10	17	5616	541.031339031339
2019	10	18	21218	664.1704213403714
2019	10	19	17944	756.7014043691485
2019	10	20	14862	741.1306015341138
2019	10	21	23082	702.249311083088



- after extracting all the month information now we can have all information in one table

```
In [39]: df_2019_agg= df_2019_04_agg.union(df_2019_05_agg)
df_2019_agg=df_2019_agg.union(df_2019_06_agg)
df_2019_agg=df_2019_agg.union(df_2019_07_agg)
df_2019_agg=df_2019_agg.union(df_2019_08_agg)
df_2019_agg=df_2019_agg.union(df_2019_09_agg)
df_2019_agg=df_2019_agg.union(df_2019_10_agg)

df_2019_agg.show(200)
```

year	month	sday	count(sday)	avg(duration_sec)
2019	4	14	9143	868.4090561084982
2019	4	15	7310	653.9667578659371
2019	4	16	13672	715.0547103569339
2019	4	17	19726	806.916607523066
2019	4	18	13505	676.8959644576083
2019	4	19	4673	618.5182965974749
2019	4	20	6604	639.3069351907934
2019	4	21	16306	1013.2635226297068
2019	4	22	21354	927.1408635384471
2019	4	23	20897	791.423505766378
2019	4	24	10624	648.4322289156627
2019	4	25	20155	743.7738526420243
2019	4	26	8179	623.4257244161878
2019	4	27	7518	605.6012237297153
2019	4	28	14192	792.2765642615558
2019	4	29	20319	757.1533540036419

- now we can do the final cleaning

rides per day	average ride time/seconds	Date
9143	868.4090561084982	2019/4/14
7310	653.9667578659371	2019/4/15
13672	715.0547103569339	2019/4/16
19726	806.916607523066	2019/4/17
13505	676.8959644576083	2019/4/18
4673	618.5182965974749	2019/4/19
6604	639.3069351907934	2019/4/20
16306	1013.2635226297068	2019/4/21
21354	927.1408635384471	2019/4/22
20897	791.423505766378	2019/4/23
10624	648.4322289156627	2019/4/24
20155	743.7738526420243	2019/4/25
8179	623.4257244161878	2019/4/26
7518	605.6012237297153	2019/4/27
14192	792.2765642615558	2019/4/28
20319	757.1533540036419	2019/4/29
23365	767.3380269634068	2019/4/30
16549	685.0921505831168	2019/5/1
19354	712.4081843546554	2019/5/2
14502	688.2785133085092	2019/5/3

only showing top 20 rows


- we can save the resulting table in Google cloud bucket that we made for our project and below we can see its partitions

The screenshot displays the Google Cloud Platform console interface. The top navigation bar shows 'Google Cloud Platform' and 'Big Data Infrastructure Pro'. The left sidebar lists 'Storage' as the active section, with sub-options like 'Browser', 'Transfer', 'Transfer Appliance', and 'Settings'. The main content area is titled 'Bucket details' for the 'infra_project' bucket. It includes tabs for 'Objects', 'Overview', 'Permissions', and 'Bucket Lock'. Below these are buttons for 'Upload files', 'Upload folder', 'Create folder', 'Manage holds', and 'Delete'. A search bar labeled 'Filter by prefix...' is present. The 'Objects' tab shows a list of files, including 'part-00002-8edbd5f8-ef43-4a25-9cb3-1e60d4...' and 'part-00003-8edbd5f8-ef43-4a25-9cb3-1e60d4...'. The file 'part-00003-8edbd5f8-ef43-4a25-9cb3-1e60d4929e99-c000.csv' is highlighted. The bottom of the screen shows the Windows taskbar with the search bar and various application icons.

Object	Size	Content Type	Storage Class	Created	Access Control	Encryption	Metadata
part-00002-8edbd5f8-ef43-4a25-9cb3-1e60d4...	35 B	application/octet-stream	Standard	11/22/19, 7:13:21 PM UTC-5	Not public	Google-managed key	-
part-00003-8edbd5f8-ef43-4a25-9cb3-1e60d4...	34 B	application/octet-stream	Standard	11/22/19, 7:13:21 PM UTC-5	Not public	Google-managed key	-
part-00003-8edbd5f8-ef43-4a25-9cb3-1e60d4929e99-c000.csv							
part-00004-8edbd5f8-ef43-4a25-9cb3-1e60d4...	35 B	application/octet-stream	Standard	11/22/19, 7:13:21 PM UTC-5	Not public	Google-managed key	-
part-00005-8edbd5f8-ef43-4a25-9cb3-1e60d4...	34 B	application/octet-stream	Standard	11/22/19, 7:13:23 PM UTC-5	Not public	Google-managed key	-
part-00006-8edbd5f8-ef43-4a25-9cb3-1e60d4...	34 B	application/octet-stream	Standard	11/22/19, 7:13:23 PM UTC-5	Not public	Google-managed key	-

Hdfs

Datanode usage histogram



In operation

Show 25 entries

Search:

Node	Http Address	Last contact	Last Block Report	Capacity	Blocks	Block pool used	Version
✓ projcet-w-0.us-central1-f.c.big-data-infrastructure-pro.internal:9866 (10.128.0.7:9866)	http://projcet-w-0.us-central1-f.c.big-data-infrastructure-pro.internal:9866	2s	80m	492.09 GB <div><div></div></div>	215	296.77 MB (0.06%)	2.9.2
✓ projcet-w-1.us-central1-f.c.big-data-infrastructure-pro.internal:9866 (10.128.0.6:9866)	http://projcet-w-1.us-central1-f.c.big-data-infrastructure-pro.internal:9866	2s	45m	492.09 GB <div><div></div></div>	215	296.31 MB (0.06%)	2.9.2

Showing 1 to 2 of 2 entries

Activate Windows
Go to Settings to activate Windows.

Previous 1 Next

- after the data is ready we can run the ARIMA time series prediction

```
In [50]: df_2019_pd=df_2019_agg.toPandas()

data = df_2019_pd['average ride time/seconds']
data.index=df_2019_pd[['Date']]

model = ARIMA(data, order=(5,1,0),dates=df_2019_pd['Date'])
model_fit = model.fit(disp=0)
print(model_fit.summary())
# plot residual errors
residuals = DataFrame(model_fit.resid)
residuals.plot()
pyplot.show()
residuals.plot(kind='kde')
pyplot.show()
print(residuals.describe())
```


ARIMA Model Results

```

=====
Dep. Variable:    D.average ride time/seconds    No. Observations:    200
Model:            ARIMA(5, 1, 0)                Log Likelihood       -1141.916
Method:           css-mle                       S.D. of innovations  72.775
Date:             Sat, 23 Nov 2019              AIC                 2297.832
Time:             01:38:02                     BIC                 2320.921
Sample:           01-15-2019                   HQIC                2307.176
                  - 01-31-2019
  
```

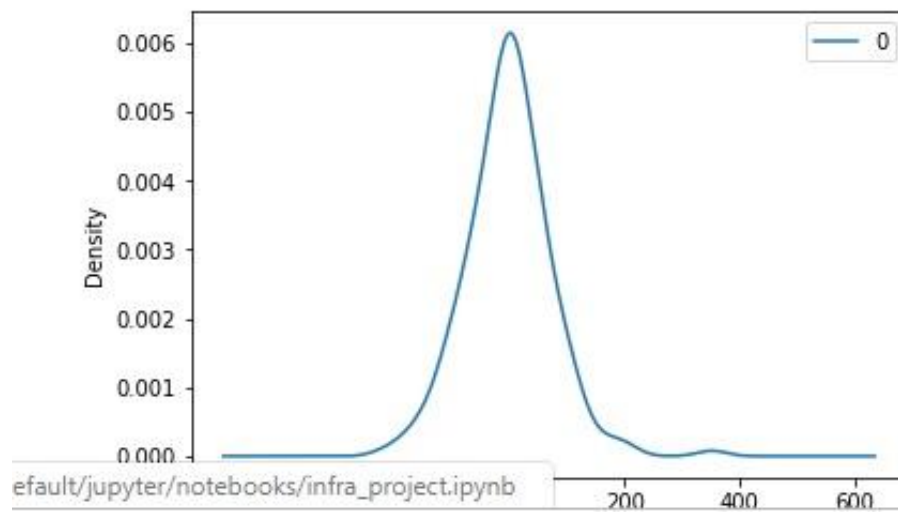

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
const                -0.4001        1.666      -0.240      0.810      -3.665        2.865
ar.L1.D.average ride time/seconds -0.4369        0.067      -6.541      0.000      -0.568      -0.306
ar.L2.D.average ride time/seconds -0.5138        0.069      -7.467      0.000      -0.649      -0.379
ar.L3.D.average ride time/seconds -0.4426        0.070      -6.309      0.000      -0.580      -0.305
ar.L4.D.average ride time/seconds -0.3492        0.069      -5.074      0.000      -0.484      -0.214
ar.L5.D.average ride time/seconds -0.3774        0.067      -5.620      0.000      -0.509      -0.246
  
```

Roots

```

=====
              Real      Imaginary      Modulus      Frequency
-----
AR.1          0.6566      -0.9596j      1.1627      -0.1545
AR.2          0.6566      +0.9596j      1.1627       0.1545
AR.3         -1.2785      -0.0000j      1.2785      -0.5000
AR.4         -0.4800      -1.1413j      1.2381      -0.3134
AR.5         -0.4800      +1.1413j      1.2381       0.3134
  
```



efault/jupyter/notebooks/infra_project.ipynb